

[Get started](#)[Log in](#)

# Cleaning up Data from Outliers


Deepika Singh

Oct 22, 2019 • 21 Min read • 24,206 Views

Oct 22, 2019 • 21 Min read • 24,206 Views

[Data](#)[Python](#)

## Introduction



- [Introduction](#)
- [Data](#)
- [Outlier Identification](#)
- [Outlier Treatment](#)
- [Conclusion](#)
- [Top](#) ^

## Introduction

The difference between a good and an average machine learning model is often its ability to clean data. One of the biggest challenges in data cleaning is the identification and treatment of outliers. In simple terms, outliers are observations that are significantly different from other data points. Even the best machine learning algorithms will underperform if outliers are not cleaned from the data because outliers can adversely affect the training process of a machine learning algorithm resulting in a loss of

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

[Disable cookies](#)[Accept cookies and close this message](#)

In this guide, you will learn about techniques for outlier identification and treatment in Python.

## Data

In this guide, we will be using a fictitious dataset of loan applications containing 600 observations and 6 variables:

- 1. Income** - Annual income of the applicant (in US dollars)
- 2. Loan\_amount** - Loan amount (in US dollars) for which the application was submitted
- 3. Term\_months** - Tenure of the loan (in months)
- 4. Credit\_score** - Whether the applicant's credit score was good ("1") or not ("0")
- 5. Age** - The applicant's age in years
- 6. Approval\_status** - Whether the loan application was approved ("1") or not ("0")

Let's start by loading the required libraries and the data.

```
4     import matplotlib.pyplot as plt
5
6     # Reading the data
7     df = pd.read_csv("data_out.csv")
8     print(df.shape)
9     print(df.info())
```

Output:

```
1         (600, 6)
2         <class 'pandas.core.frame.DataFrame'>
3         RangeIndex: 600 entries, 0 to 5
4         Data columns (total 6 columns):
5         Income                600 non-null
6         Loan_amount           600 non-null
7         Term_months            600 non-null
8         Credit_score           600 non-null
9         approval_status        600 non-null
10        Age                    600 non-null
11        dtypes: int64(6)
12        memory usage: 28.2 KB
13        None
```

The above output shows that there are 600 observations of 6 variables. All the variables have 600 records, indicating that there is no missing value in the data.

## Outlier Identification

There can be many reasons for the presence of outliers in the data.

Sometimes the outliers may be genuine,

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

[Disable cookies](#)

[Accept cookies and close this message](#)

important to understand the reasons for the outliers before cleaning them.

We will start the process of finding outliers by running the summary statistics on the variables. This is done using the ***describe()*** function below, which provides a statistical summary of all the quantitative variables.

```
python
1 df.describe()
```

Output:

count	Term_months	Credit_score	appro
667	367.10000	0.788333	0.680
98	63.40892	0.408831	0.460
00	36.00000	0.000000	0.000
000	384.00000	1.000000	0.000
000	384.00000	1.000000	1.000
000	384.00000	1.000000	1.000
000	504.00000	1.000000	1.000

Looking at the 'Age' variable, it is easy to detect outliers resulting from incorrect data. The minimum and maximum ages are 0, and 200, respectively. These are incorrect, and we will treat them later in the guide. These outliers were easy to detect, but that will not always be the case. In other cases,

techniques are discussed in the following sections.

## Identifying Outliers with Interquartile Range (IQR)

The interquartile range (IQR) is a measure of statistical dispersion and is calculated as the difference between the 75th and 25th percentiles. It is represented by the formula  **$IQR = Q3 - Q1$** . The lines of code below calculate and print the interquartile range for each of the variables in the dataset.

```
python
1     Q1 = df.quantile(0.25)
2     Q3 = df.quantile(0.75)
3     IQR = Q3 - Q1
4     print(IQR)
```

Output:

```
1     Income      3809.0
2     Loan_amount    69.5
3     Term_months      0.0
4     Credit_score      0.0
5     approval_status    1.0
6     Age            28.0
7     dtype: float64
```

The above output prints the IQR scores, which can be used to detect outliers.

The code below generates an output with the 'True' and 'False' values. Points

where the values are 'True' represent

python

```
1 print(df < (Q1 - 1.5 * IQR)) |(df >
```



Output:

		Income	Loan_amount	Term_
1				
2	0	False	False	
3	1	False	False	
4	2	False	False	
5	3	False	False	
6	4	False	False	
7	5	False	False	
8	6	False	False	
9	7	False	False	
10	8	False	False	
11	9	False	False	
12	10	False	False	
13	11	False	False	
14	12	False	False	
15	13	False	False	
16	14	False	False	
17	15	False	False	
18	16	False	False	
19	17	False	False	
20	18	False	False	
21	19	False	False	
22	20	False	False	
23	21	False	False	
24	22	False	False	
25	23	False	False	
26	24	False	False	
27	25	False	False	
28	26	False	False	
29	27	False	False	
30	28	False	False	
31	29	False	False	
32	..	...	...	
33	570	False	False	
34	571	False	False	
35	572	False	False	
36	573	False	False	
37	574	False	False	
38	575	False	False	

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

Disable cookies

Accept cookies and close this message

```

42          579    False    False
43          580    False    False
44          581    False    False
45          582    False    False
46          583    False    False
47          584    False    False
48          585    False    False
49          586    False    False
50          587    False    False
51          588    False    False
52          589    False    False
53          590    False    False
54          591    False    False
55          592    False    False
56          593    False    False
57          594    False    False
58          595    False    False
59          596    False    False
60          597    False    False
61          598    False    False
62          599    False    False
63
64          [600 rows x 6 columns]

```

## Identifying Outliers with Skewness

Several machine learning algorithms make the assumption that the data follow a normal (or Gaussian) distribution. This is easy to check with the *skewness value*, which explains the extent to which the data is normally distributed. Ideally, the skewness value should be between -1 and +1, and any major deviation from this range indicates the presence of extreme values.

The first line of code below prints the

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

Disable cookies

Accept cookies and close this message

while the second line prints the summary statistics.

```
python
1 print(df['Income'].skew())
2 df['Income'].describe()
```

Output:

```
1      6.499
2
3      count      600.000000
4      mean      7210.720000
5      std       8224.445086
6      min       200.000000
7      25%       3832.500000
8      50%       5075.000000
9      75%       7641.500000
10     max      108000.000000
11     Name: Income, dtype: float64
```

The skewness value of 6.5 shows that the variable 'Income' has a right-skewed distribution, indicating the presence of extreme higher values. The maximum 'Income' value of USD 108,000 proves this point.

## Identifying Outliers with Visualization

In the previous section, we used quantitative methods for outlier identification. This can also be achieved with visualization. Some of the common plots used for outlier detection are

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

[Disable cookies](#)


[Accept cookies and close this message](#)



The box plot is a standardized way of displaying the distribution of data based on the five-number summary (minimum, first quartile (Q1), median, third quartile (Q3), and maximum). It is often used to identify data distribution and detect outliers. The line of code below plots the box plot of the numeric variable 'Loan\_amount'.

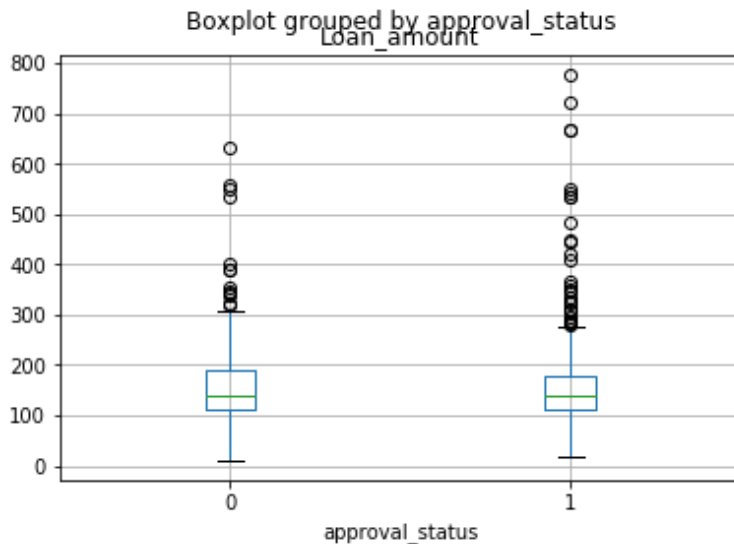
```
python
1 plt.boxplot(df["Loan_amount"])
2 plt.show()
```

Output:



In the above output, the circles indicate the outliers, and there are many. It is also possible to identify outliers using more than one variable. We can modify the above code to visualize outliers in the 'Loan\_amount' variable by the approval status.

Output:




The output shows that the number of outliers is higher for approved loan applicants (denoted by the label '1') than for rejected applicants (denoted by the label '0').

## 2. Histogram

A histogram is used to visualize the distribution of a numerical variable. An outlier will appear outside the overall pattern of distribution. The line of code below plots a histogram of the 'Income' variable, using the **hist()** function.

```
python
1 df.Income.hist()
```

Output:




The above chart shows that the distribution is right-skewed, and there are extreme higher values at the right of the histogram. This step can be repeated for other variables as well.

### 3. Scatterplot

A scatterplot visualizes the relationship between two quantitative variables. The data are displayed as a collection of points, and any points that fall outside the general clustering of the two variables may indicate outliers. The lines of code below generate a scatterplot between the variables 'Income' and 'Loan\_amount'.

```
python
1  fig, ax = plt.subplots(figsize=(12,6
2  ax.scatter(df['Income'], df['Loan_am
3  ax.set_xlabel('Income of applicants
4  ax.set_ylabel('Loan amount applied f
5  plt.show()
```



The above chart indicates that most of the data points are clustered in the lower half of the plot. The points located to the extreme right of the x-axis or the y-axis indicate outliers.

## Outlier Treatment

In the previous sections, we learned about techniques for outlier detection. However, this is only half of the task. Once we have identified the outliers, we need to treat them. There are several techniques for this, and we will discuss the most widely used ones below.

### Quantile-based Flooring and Capping

In this technique, we will do the flooring (e.g., the 10th percentile) for the lower values and capping (e.g., the 90th percentile) for the higher values. The lines of code below print the 10th and 90th percentiles of the variable


We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

Disable cookies

Accept cookies and close this message

be used for quantile-based flooring and capping.

```
python
1 print(df['Income'].quantile(0.10))
2 print(df['Income'].quantile(0.90))
```




Output:

```
1      2960.1
2     12681.0
```

Now we will remove the outliers, as shown in the lines of code below. Finally, we calculate the skewness value again, which comes out much better now.

```
python
1 df["Income"] = np.where(df["Income"]
2 df["Income"] = np.where(df["Income"]
3 print(df['Income'].skew())
```



Output:

```
1      1.04
```

## Trimming

In this method, we completely remove data points that are outliers. Consider the 'Age' variable, which had a minimum value of 0 and a maximum value of 200.

The first line of code below creates an

line drops these index rows from the data, while the third line of code prints summary statistics for the variable.

After trimming, the number of observations is reduced from 600 to 594, and the minimum and maximum values are much more acceptable.

```
python
1 index = df[(df['Age'] >= 100)|(df['A
2 df.drop(index, inplace=True)
3 df['Age'].describe()
```



Output:

1	count	594.000000
2	mean	50.606061
3	std	16.266324
4	min	22.000000
5	25%	36.000000
6	50%	50.500000
7	75%	64.000000
8	max	80.000000
9	Name: Age, dtype: float64	

## IQR Score

This technique uses the IQR scores calculated earlier to remove outliers. The rule of thumb is that anything not in the range of **(Q1 - 1.5 IQR) and (Q3 + 1.5 IQR)** is an outlier, and can be removed. The first line of code below removes outliers based on the IQR range and

shape of this data, which comes out to be 375 observations of 6 variables. This shows that for our data, a lot of records get deleted if we use the IQR method.

```
python
1 df_out = df[~((df < (Q1 - 1.5 * IQR)
2 print(df_out.shape)
```



Output:

```
1 (375, 6)
```

## Log Transformation

Transformation of the skewed variables may also help correct the distribution of the variables. These could be logarithmic, square root, or square transformations. The most common is the logarithmic transformation, which is done on the 'Loan\_amount' variable in the first line of code below. The second and third lines of code print the skewness value before and after the transformation.

```
python
1 df["Log_Loanamt"] = df["Loan_amount"]
2 print(df['Loan_amount'].skew())
3 print(df['Log_Loanamt'].skew())
```



Output:

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

[Disable cookies](#)

[Accept cookies and close this message](#)

```
2.8146019248106815  
-0.17792641310111373  
1  
2
```

The above output shows that the skewness value came down from 2.8 to -0.18, confirming that the distribution has been treated for extreme values.

## Replacing Outliers with Median Values

In this technique, we replace the extreme values with median values. It is advised to not use mean values as they are affected by outliers. The first line of code below prints the 50th percentile value, or the median, which comes out to be 140. The second line prints the 95th percentile value, which comes out to be around 326. The third line of code below replaces all those values in the 'Loan\_amount' variable, which are greater than the 95th percentile, with the median value. Finally, the fourth line prints summary statistics after all these techniques have been employed for outlier treatment.

```
python  
1 print(df['Loan_amount'].quantile(0.5)  
2 print(df['Loan_amount'].quantile(0.9  
3 df['Loan_amount'] = np.where(df['Loa  
4 df.describe()
```

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

Disable cookies

Accept cookies and close this message




## Output:

```

1          140.0
2          325.75000000000001
3
4
5          |          | Income          | Loan_
6          | - - - - - | - - - - - | - - - - -
7          | count     | 594.000000 | 594.0
8          | mean      | 6112.375421 | 144.2
9          | std       | 3044.257269 | 53.03
10         | min       | 2960.000000 | 10.00
11         | 25%      | 3831.500000 | 111.0
12         | 50%      | 5050.000000 | 140.0
13         | 75%      | 7629.000000 | 171.0
14         | max      | 12681.000000 | 324.0

```



## Conclusion

In this guide, you have learned methods of identifying outliers using both quantitative and visualization techniques. You have also learned techniques for treating the identified outliers. Your usage of these techniques will depend on the data, the problem statement, and the machine learning algorithm selected for building the model.

To learn more about data preparation and building machine learning models using Python, please refer to the

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

[Disable cookies](#)

[Accept cookies and close this message](#)


2. [Linear, Lasso, and Ridge Regression with scikit-learn](#)
3. [Non-Linear Regression Trees with scikit-learn](#)
4. [Machine Learning with Neural Networks Using scikit-learn](#)
5. [Validating Machine Learning Models with scikit-learn](#)
6. [Ensemble Modeling with scikit-learn](#)
7. [Preparing Data for Modeling with scikit-learn](#)
8. [Interpreting Data Using Descriptive Statistics with Python](#)

To learn more about building deep learning models using **Keras**, please refer to the following guides:

1. [Regression with Keras](#)
2. [Classification with Keras](#)



LEARN MORE



<https://www.pluralsight.com/offer/2020/august-free-weekend>  
**IS HERE! FRIDAY 8/14 - SUNDAY 8/16**

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

Disable cookies

Accept cookies and close this message

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here \(/privacy\)](#).

SOLUTIONS

ACCEPT COOKIES AND CLOSE THIS MESSAGE

[Disable cookies](#)

Pluralsight Skills (/product/skills)

We use cookies to make interactions with our websites and services easy and meaningful. For more information about the cookies we use or to find out how you can disable cookies, [click here](#).

[Disable cookies](#)

Accept cookies and close this message